Paired differences show substantial uncertainty

Each method minus Single RAG; exploratory 95% question-bootstrap intervals

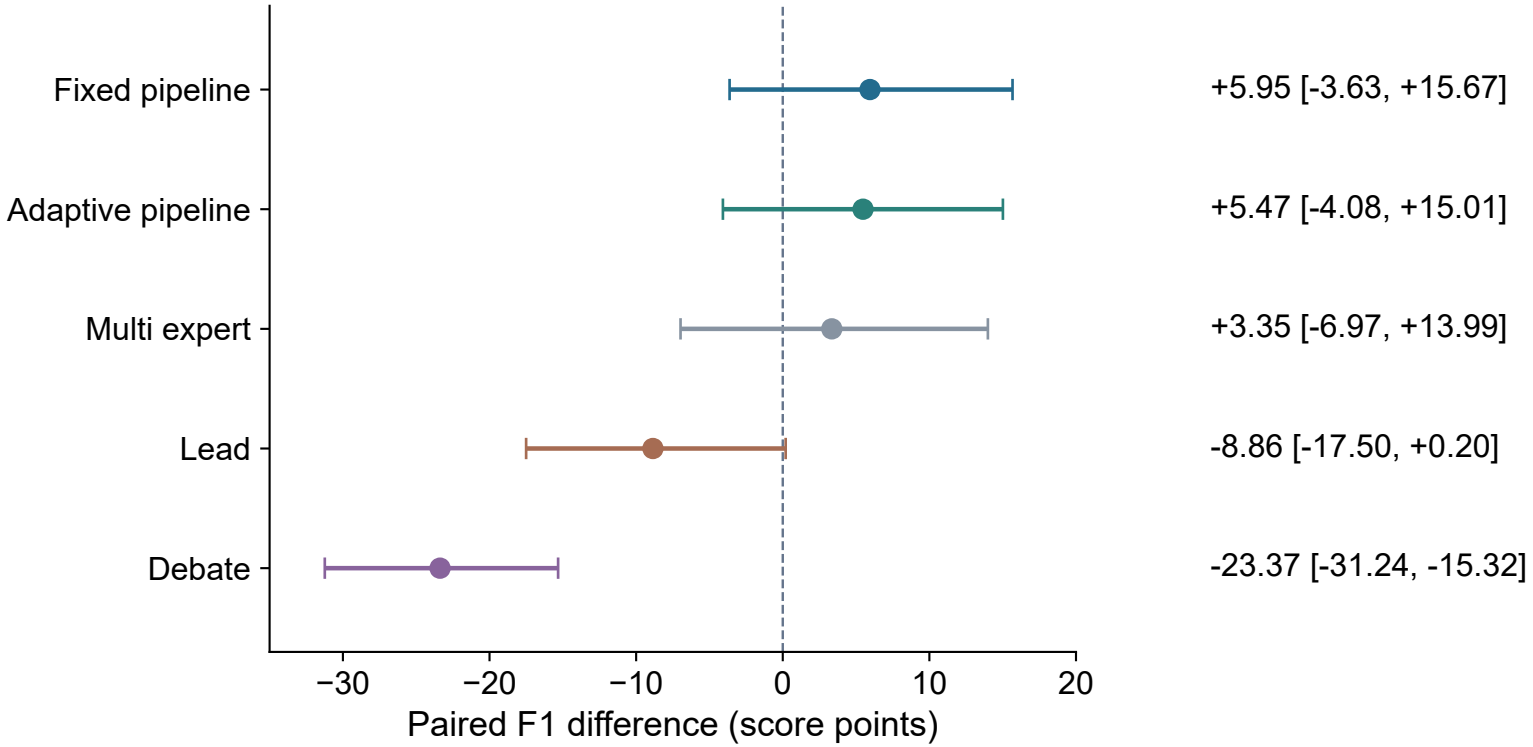

n = 100 paired questions; 2,000 bootstrap resamples; one model run per method.

Intervals are unadjusted and describe question resampling, not model-run variability.